

Piuma's Manual

A refactoring library for Scala compiler extension

Amanj Sherwany

June 26, 2015

1 Abstract

Compiler plugins enable Scala to be extended with new functionality by adding compiler passes that perform additional static checking, code generation, or code transformations. However, they are often difficult to build. A plugin can perform arbitrary code transformations, easily allowing a developer to generate incorrect code. Moreover, Scala's compiler assumes many complex, sometimes undocumented invariants, requiring plugin developers to acquire intimate knowledge of the design and implementation of the compiler.

To address these issues we introduce Piuma. Piuma is a library that provides, first, an API to perform many common refactoring tasks needed by plugin writers, and, second, a DSL to eliminate much of the boilerplate code required for plugin development.

2 Introduction

Scala provides an API to extend its compiler. Developers write *compiler plugins* to add new passes to the base Scala compiler. Compiler plugins are a useful tool in DSL development because they allow passes to, for instance, perform additional static analysis, to instrument code with additional dynamic checks, to perform optimizations, or to generate code for new language features.

Plugin developers can make nearly arbitrary changes to the base compiler, permitting implementation of complex language features, but unfortunately also permitting plugins to violate invariants assumed by the compiler. Breaking these often undocumented invariants may cause the compiler to generate incorrect Java bytecode or even to crash.

There are many ways to generate malformed bytecode using Scala compiler plugins. For example, a compiler plugin can add a “ghost” field to a class that can be seen by the Java VM when running the code, but not by the Scala compiler itself when importing the generated bytecode. This problem occurs because the Scala compiler embeds Scala-specific type information into the generated Java bytecode. If a plugin were to add a field but omit this extra information, another instance of the Scala compiler would not even see the field even though

```

// define a new extension
@plugin(Phase1)
class Example

// define a compilation phase
@treeTransformer("example")
class Phase1 {
  rightAfter("parser")
  before(List("patmat"))

  def transform(tree: Tree): Tree = super.transform(tree)
}

```

Figure 1: Using Piuma DSL for creating a simple compiler extension

the field is present in the Java bytecode. This can result in other compilation units failing to compile correctly.

Since plugins add passes to the Scala compiler, running a plugin at the wrong point in the compilation can also allow bad code to be generated. For instance, a plugin could rename a field to have the same name as an existing field. If the plugin did this after the Scala type-checker runs, the error would not be detected and the bytecode would contain a duplicate field.

Piuma is a refactoring library for the Scala compiler for more easily implementing correct compiler plugins. The library provides a set of refactoring methods to, for instance, safely rename definitions, to add new class, trait, or object members, or to extract expressions into methods. The library also provides a DSL for generating the boilerplate code necessary for writing Scala compiler plugins. In the following section we show how Piuma is used.

2.1 The Piuma library

Suppose a developer is writing a plugin that performs partial evaluation. To do this, she needs to introduce specialized versions of methods into a class’s companion object.¹ If the companion object does not exist, the plugin must introduce one. This can be a tedious task, as shown in Figure 2.

Every AST node in Scala has a `Symbol` attached to it, which includes type and other information about the node. The code first introduces a symbol for the companion object, called a “module” in the Scala compiler API. The code then initializes the symbol with its owner (i.e., its containing package or class), its type, and its parents (i.e., its supertypes). It then updates the owner of the method we want to insert into the module (code elided). It introduces a default constructor that calls the constructor of the object’s superclass, `AnyRef`, and appends the new constructor to the body of the module. Finally, it types the

¹In Scala, a companion object for a class is singleton object with the same name as the class that contains “static” members for the class.

```

// clazz is the symbol of the class that we want to
// have its companion object
// mthd is a method we want to include in the module
val moduleName = clazz.name
val owner = clazz.owner
val moduleSymbol = clazz.newModule(moduleName.toTermName,
                                   clazz.pos.focus, Flags.MODULE)
val moduleClassSymbol = moduleSymbol.moduleClass
moduleSymbol.owner = owner
moduleClassSymbol.owner = owner

val parents = List(Ident(definitions.AnyRefClass))
val moduleType = ClassInfoType(parents.map(_.symbol.tpe),
                                newScope, moduleClassSymbol)
moduleClassSymbol setInfo moduleType
moduleSymbol setInfoAndEnter moduleClassSymbol.tpe

// elided code: plugin writer needs to fix the owner
// of mthd and its children

val constSymbol =
  moduleClassSymbol.newClassConstructor(moduleSymbol.pos.focus)

constSymbol.setInfoAndEnter(MethodType(Nil, moduleSymbol.info)

val superCall = Select(Super(This(tpnme.EMPTY), tpnme.EMPTY),
                       nme.CONSTRUCTOR)
val rhs = Block(List(Apply(superCall, Nil)),
                 Literal(Constant(())))
val constructor = DefDef(constSymbol, List(Nil), rhs)

localTyper.typed {
  ModuleDef(moduleSymbol, Template(parents, noSelfType,
                                   List(cnstrct, mthd)))
}

```

Figure 2: This listing shows how a new companion object is created for a class using the Scala compiler API.

```

// clazz is the symbol of the class for which we
// want its companion object
// mthd is a method we want to include in the module
val module0 = clazz.mkCompanionObject
val module = module0.addMember(mthd)

```

Figure 3: This listing shows how a new companion object is created for a class using the Piuma DSL.

module tree. Failing to do any of these steps may lead to an exception during compilation, or worse, the compiler might silently generate malformed bytecode. With Piuma we can introduce a companion object in just a few lines of code, as shown in Figure ?? . In Section ?? we describe the design of the Piuma utilities.

3 Piuma and SBT

SBT is an open source build tool for Scala, Java and more. It is similar to Java's Maven or Ant. It is mainly used for Scala, therefore we show how Piuma is used with SBT.

3.1 Piuma as a dependency

To use Piuma you need to first add it to the list of the *resolvers* as follows:

```
resolvers += "Piuma's Repo" at
             "http://inf.usi.ch/phd/sherwany/repos/piuma"
```

Then, you need to add Piuma in the project dependencies list:

```
libraryDependencies +=
  Seq("ch.usi.inf.l3" %% "piuma" % "0.1-SNAPSHOT"),
```

Since Piuma relies on Scala Macros, you need to add macros to the compiler plugins list, please note that you need to run Scala 2.11 or above.

```
addCompilerPlugin("org.scalamacros" % "paradise" %
                  paradiseVersion cross CrossVersion.full)
```

3.2 Deploying Plugins

Because of the limitations of scalac, plugin jar files can only be fat jars, therefore you need to include all the dependencies in your plugin's jar file. This can be easily done using an SBT plugin like `sbt-assembly`, which can be found here: <https://github.com/sbt/sbt-assembly>. For more information please consult its help. To deploy fat jar files, add this to your settings list:

```
settings(
  artifact in (Compile, assembly) ~= { art =>
    art.copy('classifier' = Some("assembly"))
  }
) settings (assemblySettings ++ addArtifact(artifact in (
  Compile, assembly),
  assembly))
```

3.3 A Complete SBT Example

```
lazy val plugin = Project(
  id    = "plugin",
  base = file("plugin")
) settings (
  // general settings
  scalaVersion := "2.11.1",
  resolvers += "Piuma's Repo" at
    "http://inf.usi.ch/phd/sherwany/repos/piuma",
  resolvers += Resolver.sonatypeRepo("snapshots"),
  resolvers += Resolver.sonatypeRepo("releases"),

  publishMavenStyle := true,
  resourceDirectory in Compile <= baseDirectory(_ / "resources"),
  publishTo := Some(Resolver.sftp("Plugin's Repo",
    "some.url.to.repo",
    "path/to/repo")),

  // Supposing that ``scalac-plugin.xml`` is installed in the directory
  // resources under the plugin directory
  resourceDirectory in Compile <= baseDirectory(_ / "resources"),

  libraryDependencies ++=
    Seq("ch.usi.inf.l3" %% "piuma" % "0.1-SNAPSHOT"),

  addCompilerPlugin("org.scalamacros" % "paradise" %
    paradiseVersion cross CrossVersion.full)
) settings (
  artifact in (Compile, assembly) ~= { art =>
    art.copy('classifier' = Some("assembly"))
  }
) settings (assemblySettings ++
  addArtifact(artifact in (Compile, assembly),
    assembly) : _*) dependsOn(library)
```

To use the resulted plugin, in another project it is as simple as having the following SBT build script:

```
lazy val someName = Project (
  id    = "someName",
  base = file("someName")
) settings (
  // general settings
  scalaVersion := "2.11.1",
  resolvers += "Piuma's Repo" at
    "http://inf.usi.ch/phd/sherwany/repos/piuma",
```

```

    resolvers += Resolver.sonatypeRepo("snapshots"),
    resolvers += Resolver.sonatypeRepo("releases"),

    autoCompilerPlugins := true,
    addCompilerPlugin(
      "ch.usi.inf.l3" %% "plugin" %
      "0.1-SNAPSHOT" classifier "assembly" changing())
  )

```

4 The Piuma DSL

The Piuma DSL extends Scala with features that facilitate defining compiler plugins and their components. In this section, we explain the design and use of this DSL in detail. We start by describing the general structure of a Scala compiler plugin and then describe the DSL constructs and the functionality they provide. Finally, we briefly describe how the DSL is implemented.

The Scala compiler consists of a sequence of compilation phases. Developers extend the compiler by creating plugins, composed of one or more phases inserted into this sequence. Compiler plugins are implemented by extending the `Plugin` class and providing a list of `PluginComponent`. Each of these components specifies a compiler phase and where it occurs in the compilation sequence. They also provide factory methods for creating the tree and symbol transformers that implement the phase.

The Piuma DSL extends Scala with four class annotations: `@plugin`, `@checker`, `@treeTransformer`, and `@infoTransformer`. The `@plugin` annotation generates boilerplate code for a compiler plugin itself. The other annotations generate boilerplate code for different `PluginComponent` implementations: `@checker` generates a type-checking component, `@treeTransformer` generates an AST-transforming component, and `@infoTransformer` generates a symbol-transforming component. Since the Scala compiler requires plugins and phases to be concrete classes, these annotations cannot appear on traits, abstract classes, or singleton object. Annotated classes can still implement other traits. Figure 4 shows a sketch of the syntax of a compiler extension in the DSL.

The Piuma DSL is implemented using Scala's *annotation macros*. For each annotation, macro expansion modifies the annotated class to extend a corresponding class from the Scala compiler API. It then mixes-in appropriate Piuma traits to facilitate access to the Piuma library.

4.1 The `@plugin` annotation

The `@plugin` annotation is placed on a class that implements a compiler plugin. After macro expansion, the annotated class automatically extends the Scala class `Plugin`. The list of components provided by the plugin are specified as annotation arguments. Optionally, the class may provide a short description

```

@plugin(MyChecker, MyTransformer, MyInfoTransformer)
class MyPlugin {
  describe("short_description")
  ...
}

@checker("my_checker")
class MyChecker {
  plugin MyPlugin // optional

  after(List("phase1", "phase2", ...)) // optional
  rightAfter("phase1") // optional
  before(List("phase1", "phase2", ...)) // optional

  def check(unit: CompilationUnit): Unit = ...
  ...
}

@treeTransformer("my_transformer")
class MyTransformer {
  plugin MyPlugin // optional

  after(List("phase1", "phase2", ...)) // optional
  rightAfter("phase1") // optional
  before(List("phase1", "phase2", ...)) // optional

  def transform(tree: Tree): Tree = ...
  ...
}

@infoTransformer("my_info_transformer")
class MyInfoTransformer {
  plugin MyPlugin // optional

  after(List("phase1", "phase2", ...)) // optional
  rightAfter("phase1") // optional
  before(List("phase1", "phase2", ...)) // optional

  def transform(tree: Tree): Tree = ...
  def transformInfo(sym: Symbol, tpe: Type): Type = ...
  ...
}

```

Figure 4: Syntax of the Piuma DSL

of its purpose, used when generating command-line usage information for the compiler.

4.2 Component annotations

The three annotations `@checker`, `@treeTransformer`, and `@infoTransformer` are used to annotate classes that implement plugin components. Macro expansion generates boilerplate code in the annotated class for inserting the phase into the execution order. These annotations also specify the name of the phase using an annotation parameter.

In the body of a class annotated with one of the phase annotations, the programmer can optionally specify the class of the compiler plugin itself using the syntax `plugin PluginClass`. This introduces a field of the appropriate type into the class that refers to the plugin object. This field can be used to share information across the plugin's various compiler phases.

`@checker` .

This annotation is placed on classes that implement type-checking phases. A checker phase cannot perform AST transformations but can perform static analysis of a compilation unit. The class must implement a method with the signature: `check(CompilationUnit): Unit`. After macro expansion, the class extends the `Plugin-Component` class from the Scala compiler API and implements a factory method for creating compiler phase objects that invoke the `check` method for each compilation unit.

`@treeTransformer` .

The `@treeTransformer` annotation is used for component classes that implement AST transformations. Annotated classes must implement a method with the signature: `transform(Tree): Tree`. After expansion, the class extends `Plugin-Component` **with** `Typing-Transform`. The expanded class creates a `Tree-Transformer` that traverses the AST and invokes the provided `transform` method at each node.

`@infoTransformer` .

The last annotation is `@infoTransformer`, which is placed on classes that transform types in the AST. The annotation is similar to `@treeTransformer`; however, classes must provide not only a `transform(Tree): Tree` method, but also a `transformInfo(Symbol, Type): Type` method. After expansion, an annotated class will extend `Plugin-Component` **with** `Info-Transform`. A generated `Info-Transformer` class traverses the AST and invokes the provided `transform` and `transformInfo` methods for each node and symbol encountered.

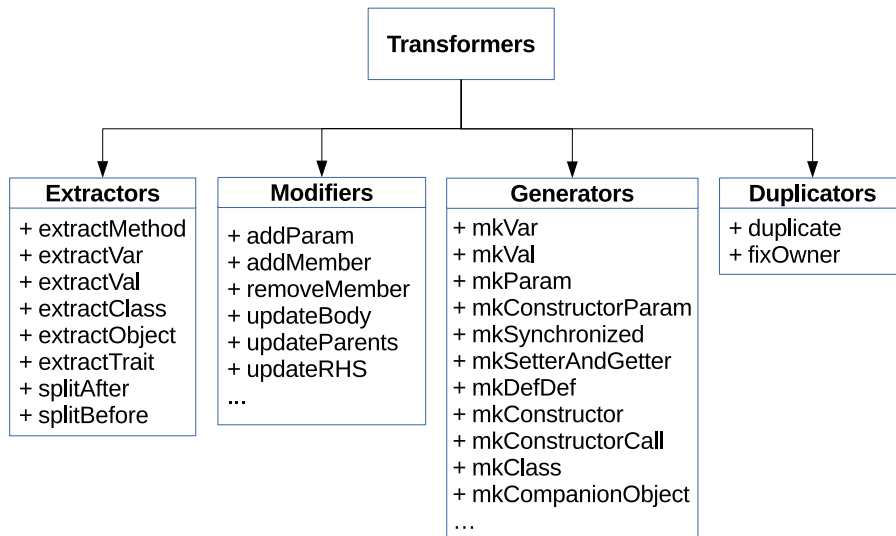


Figure 5: The high-level design of Piuma's Library

5 The Piuma library

Piuma offers a rich set of utilities for generating and refactoring Scala compiler ASTs. In general, these methods implement common refactoring and creation patterns when writing compiler plugins. They are implemented using the AST generators of the Scala compiler API and Scala's quasiquotes. We aimed to provide a library that is easy to use, yet flexible and expressive. Library users can use the Scala compiler API alongside library code if they need lower-level access to the compiler internals.

In this section, we discuss the design of the refactoring library and demonstrate its use through use cases and code examples. The library is divided into four main categories, as shown in Figure 5. The reader can refer to Piuma's project page² for documentation and more examples.

5.1 Tree extractors

Methods in the *tree extractor* category permit the selection of a sequence of ASTs to be placed in another, compatible class, object, trait, method or variable. Since Scala ASTs are immutable, extractor operations generate new ASTs and do not affect the original tree to which they are applied.

Figure 6(a) shows a code snippet that demonstrates a refactoring using a pair of tree extractor functions. Figures 6(b) and (c) show the effect of applying the

²<https://github.com/amanjpro/piuma>

refactoring to a small method. The extractor function `splitAfter` scans the AST of the body of the original method `foo` in Figure 6(b). When `splitAfter` finds the first occurrence of an AST node that satisfies a given predicate and returns the AST before and include the node and the AST after the node. In the example, the predicate matches a call to the `splitMe` method, specified using the Scala quasiquotes library. Thus `splitAfter` returns the AST for the body of `foo` up to and including the call to `splitMe`, and another AST for the rest of the body.

It is possible to insert either or both of the ASTs returned by `splitAfter` into ASTs. In Figure 6(c), `suffix` becomes the body of a new method `bar`. This is done by calling `extractMethod`, which takes a list of trees and a new method name as parameters. The new name must be unique in its scope, otherwise an exception is thrown.

The `extractMethod` operation handles free variables in the given method's body by adding parameters to the extracted method. That is, all free variables in `suffix` become parameters of `bar`. `extractMethod` also creates the AST of the new method and types it, and generates a call to the extracted method that can be substituted into the original method, as described in the next section. Piuma also creates a symbol for the new method, rebinding symbols of the extracted trees to their new owner, and placing them in the symbol table of the method owner. Without Piuma, the programmer would have to perform all these operations manually using the Scala compiler API.

Other useful Piuma extractors allow member extraction. It is possible, for instance, to extract an inner class and make it an outer class (using `extractClass`), to convert a local variable to a field (using `extractVar` or `extractVal`), or to move fields or methods across classes.

Applying Piuma's extractor functions can sometimes yield malformed trees. For example, if we add a block extracted from a method that contains a return statement to a class, the class will not type check. We rely on Scala's type checker to report such errors to the programmer.

5.2 Tree modifiers

Tree modifier utilities change AST nodes, for instance, adding or removing method parameters (`addParam` and `removeParam`, respectively), modifying a class or trait's body (`updateBody`, `addMember`, and `removeMember`) or its supertypes (`updateParents`). The `rename` method is used to change the name of a class, method, or variable. When a field is renamed, Piuma handles renaming its setters and getters as well. The tree modifiers also support updating method bodies (`updateRHS`). As an example, we can modify method `foo` from Figure 6(b) to call the extracted method, as shown in Figure 6(c). We do this with the call to `updateRHS`, as shown in Figure 6(a).

```

// orig is the original method
// body is the body of orig

// Split body after the call to splitMe(),
val (prefix, suffix) = body.splitAfter((x:Tree) => {
  x == q"splitMe()"
})

// Extract the code after the split into method bar
val (extracted, apply) = extractMethod(
  suffix,      // The body of the extracted method
  "bar"        // the name of the extracted method
  orig.symbol // the symbol of the original method
).get

// Replace the extracted code with a call to the new method
orig.updateRHS(Block(prefix, apply))

```

(a) A short example of tree extractor and tree modifier utilities

```

def foo(a: Int, b: String,
        c: Int): Int = {

  println(a)
  val d = c
  splitMe()
  println(d)
  a + b.size
}

```

(b) Before applying the refactoring in
(a)

```

def foo(a: Int, b: String,
        c: Int): Int = {
  println(a)
  val d = c
  splitMe()
  bar(a, b, d)
}
def bar(a: Int, b: Int,
        d: Int): Int = {
  println(d)
  a + b.size
}

```

(c) After applying the refactoring in (a)

Figure 6: A short example of a Piuma refactoring

5.3 Tree duplicators

Tree duplicators are mainly used to implement other Piuma functionality, but can also be useful when writing plugins. For example, if a programmer needs to specialize a method, they can duplicate it first, then remove the specialized parameter from the duplicated tree, and substitute it with a value, as shown in Figure 7.

Piuma's `duplicate` method differs from the compiler API's `duplicate` method in that it handles the creation of a new symbol for the new tree, and

```
// orig is the method to be specialized  
// p is the parameter we want to specialize  
// v is the value that we want to specialize p with  
val specialized = orig.duplicate("new_name").removeParam(p, v)
```

Figure 7: A simple example showing how a tree duplicator works

changes the binding of the node's children's symbols from their original owner to the duplicate.

The `fixOwner` method traverses an AST that has been duplicated and inserted into the AST at another location. It rebinds symbols in the traversed AST to change their owner (i.e., containing method, class, etc.) to the symbol at the new location in the AST.

5.4 Tree generators

Piuma offers various AST generators that are simpler to use than the Scala compiler API tree generators. For instance, they handle setting the required flags that distinguish trees used to represent more than one syntactic construct (e.g., `var` and `val`, or class and trait trees). Symbol creation and binding is also handled for a generated AST, including all its descendant ASTs. Piuma generators facilitate other tasks that require multiple setup steps, such as surrounding the body of a method with a synchronization block, creating constructor parameters and constructor calls, and others.

6 Getting More Help

Piuma's library is fully documented, you are encouraged to visit the project's repository and clone it, play with it and understand it. The project's repository can be found at <https://github.com/amanjpro/piuma>. You can generate the Scaladoc by running

```
sbt piuma/api
```

You can also take a look at the many plugins that use Piuma and included in the project, namely (Avro, Kara, Mina, Atomic Scala and ScalaDyno). You are welcome to report bugs, suggest improvements and/or contribute to the project.

7 Software Requirements

Piuma should have no problem with Scala 2.11, and in theory it should work with newer Scala versions.

8 Licence

Copyright (c) <2013-2014>, Amanj Sherwany <<http://www.amanj.me>>
All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- * Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- * Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
- * Neither the name of the <organization> nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL <COPYRIGHT HOLDER> BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.